OPEN GOVERNANCE STANDARD · V1.0

Govern Your AI Agents Before They Govern You

The Sentinel Framework is an open standard for deploying AI agents safely, accountably, and auditably — built for the era when AI doesn't just answer questions, it takes action.

Read the Standard | Compliance Checklist

> *"Don't give AI a job. Give it a task. Know every token it spends. Put a human in the middle — not at the end."*
>
> — Christopher Gyening, Architect · Sentinel Framework v1.0

The Problem with AI Today

Organizations are giving AI jobs — broad, vague mandates with wide access and no accountability trail. When something goes wrong, no one can explain what happened, who authorized it, or how to prevent it next time. Sentinel fixes this.

THE FIVE PILLARS

A Defense-in-Depth Architecture for Agentic AI

Each pillar addresses a specific failure mode. Together, they form a complete governance layer.

PILLAR 01

Token Accountability

Every token generated is tied to a Budget Identity — giving you a real-time audit trail of how the AI thought, not just what it said.

"Your AI has an expense report. Audit it."

Read more →

PILLAR 02

Human-In-The-Middle

Move beyond passive "human-in-the-loop" to active intervention. Humans intercept, correct, and steer agents before high-stakes actions execute.

"Your CFO doesn't approve budgets after the money's spent."

Read more →

PILLAR 03

RAG vs. LAG

RAG grounds the AI in truth. LAG monitors its trajectory over time. When drift exceeds a threshold, the session is killed before the action phase.

"GPS + black box. One navigates. One proves what happened."

Read more →

PILLAR 04

Least Privilege Agency

Every agent receives an Ephemeral Scoped Role — a time-limited, task-specific permission token. It expires when the task ends. No lingering access.

"Contractors get key cards, not the building access code."

Read more →

PILLAR 05

Tasks Over Jobs

AI shouldn't have jobs — it should have tasks. Atomic, scoped, stateless. Once complete, memory is wiped. No long-term goal drift possible.

"You hire a plumber to fix a pipe. Not to run your building."

Read more →

IMPLEMENTATION

Three Phases to Compliance

You don't need to implement everything at once. The Sentinel maturity model lets you progress at your own pace.

1. **Phase 1 — Foundation (Weeks 1–4)**
   Audit your AI tools. Classify Jobs vs. Tasks. Remove unused permissions. Add basic logging. Define your first task specifications.

2. **Phase 2 — Governance Layer (Months 2–3)**
   Deploy token monitoring, HITM gates, and scoped permissions. Every agent action is logged. High-risk actions require human sign-off.

3

Phase 3 — Sentinel Standard (Months 4–6)

Full RAG/LAG drift monitoring, stateless agent architecture, and compliance certification. You can answer any question about any AI decision in the last 90 days.

WHO IT'S FOR

Built for Everyone Who Deploys AI

Whether you're a board member, an engineer, or a compliance lead — Sentinel gives you a common language.

Executives

Understand what your AI agents are doing, what they have access to, and who is accountable.

Engineers

A concrete architecture for building agentic systems that are auditable by design.

Legal & Compliance

A framework your team can audit, cite, and enforce — not a vague policy document.

Researchers

An open, versioned standard for studying AI governance in production environments.

Start Building Sentinel-Compliant AI

The standard is open source. Read it, use it, build on it — just give credit where it's due.

View on GitHub

Take the Checklist

```
[![Sentinel Compliant](https://img.shields.io/badge/Sentinel-Compliant-blue)]
(https://github.com/gyeningcorp/sentinel-framework)
```

© 2026 Christopher Gyening / gyeningcorp · The Sentinel Framework is licensed under Apache 2.0 · GitHub